

BIOL 343

Applied Bioinformatics

Counting

Dr. Nic Wheeler

Counting converts mapped reads to expression values


Expression values will be used by statistical models to identify DEGs

- Recall the goal of our RNA-seq experiments...
 - Treatment vs Control
 - Mutant vs Wild type
 - *Identify differentially expressed genes (DEGs)*
- DEGs will be identified using statistical tests comparing **expression values** of transcripts/genes
- Expression values will be calculated based on the number of reads that **align/map** to a given genomic locus
- There are different ways to count reads

Counting converts mapped reads to expression values

Expression values will be used by statistical models to identify DEGs

Sample 1



9 reads per million

Sample 2



18 reads per million


RNA

Library


Reads

Expression values

Differentially expressed gene?



Higher expression in CTRL
than PZQ...do we see that?



How are reads converted to expression values?

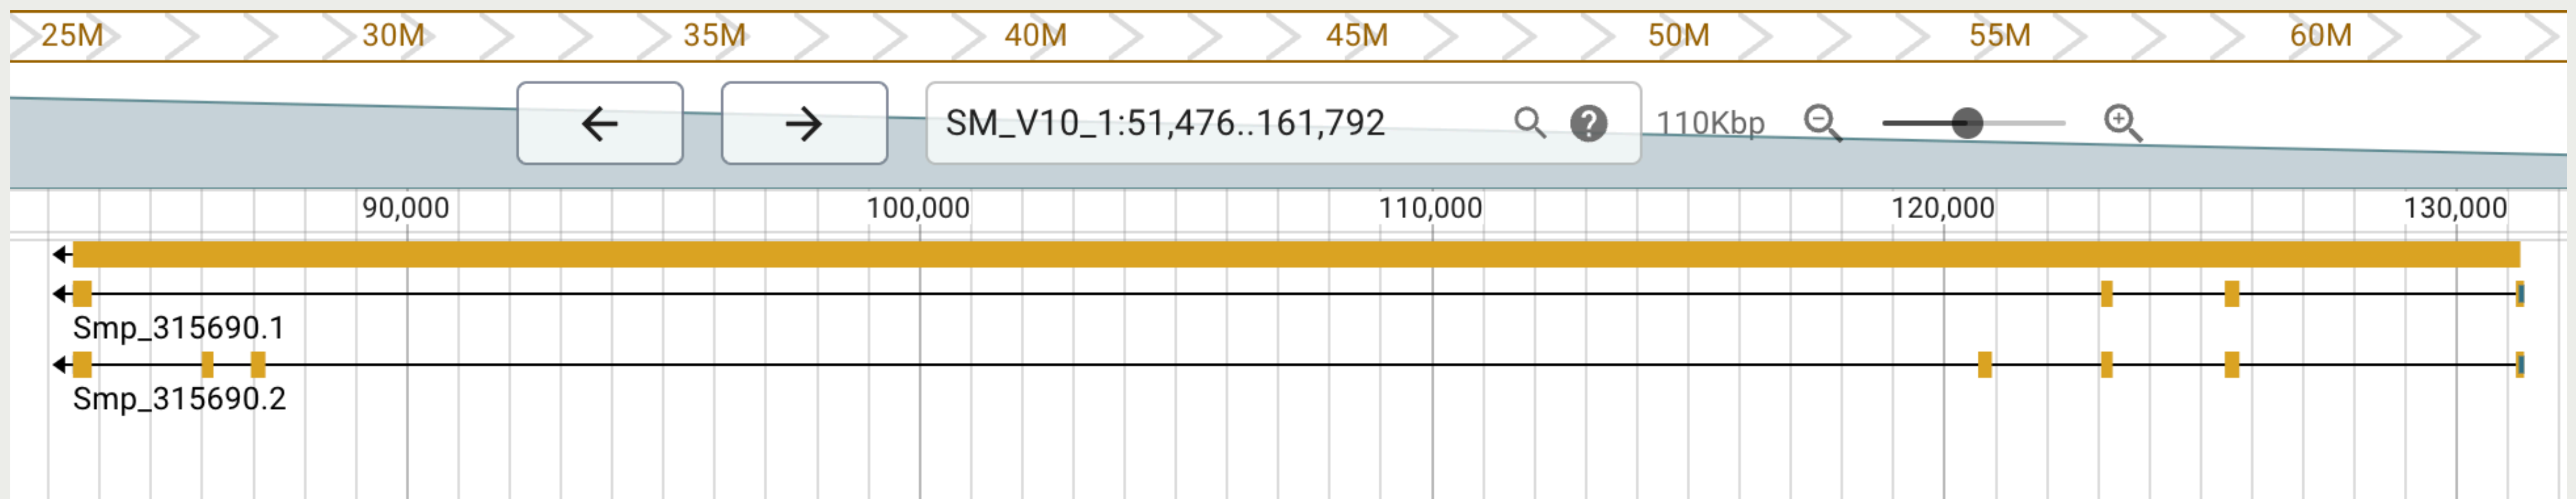
Many different situations must be considered

- What if a read only overlaps 1 nt of the gene?
- What if a read overlaps two genes?
- What if a gene has a lot of reads at one location and no reads at a different location?
- What if a read overlaps mostly intron?
- How do you link reads to different isoforms of the same gene?
- How do you normalize for genes that have variable lengths?

The featureCounts approach

Features and meta-features


- Consider SM_V10_1:80,002..135,160
 - Smp_315690 (gene)
 - Smp_315690.1 and Smp_315690.2 (transcripts)
 - Differ by the 3 exons (found in .2 but not .1)



[illegible]

The featureCounts approach

Features and meta-features



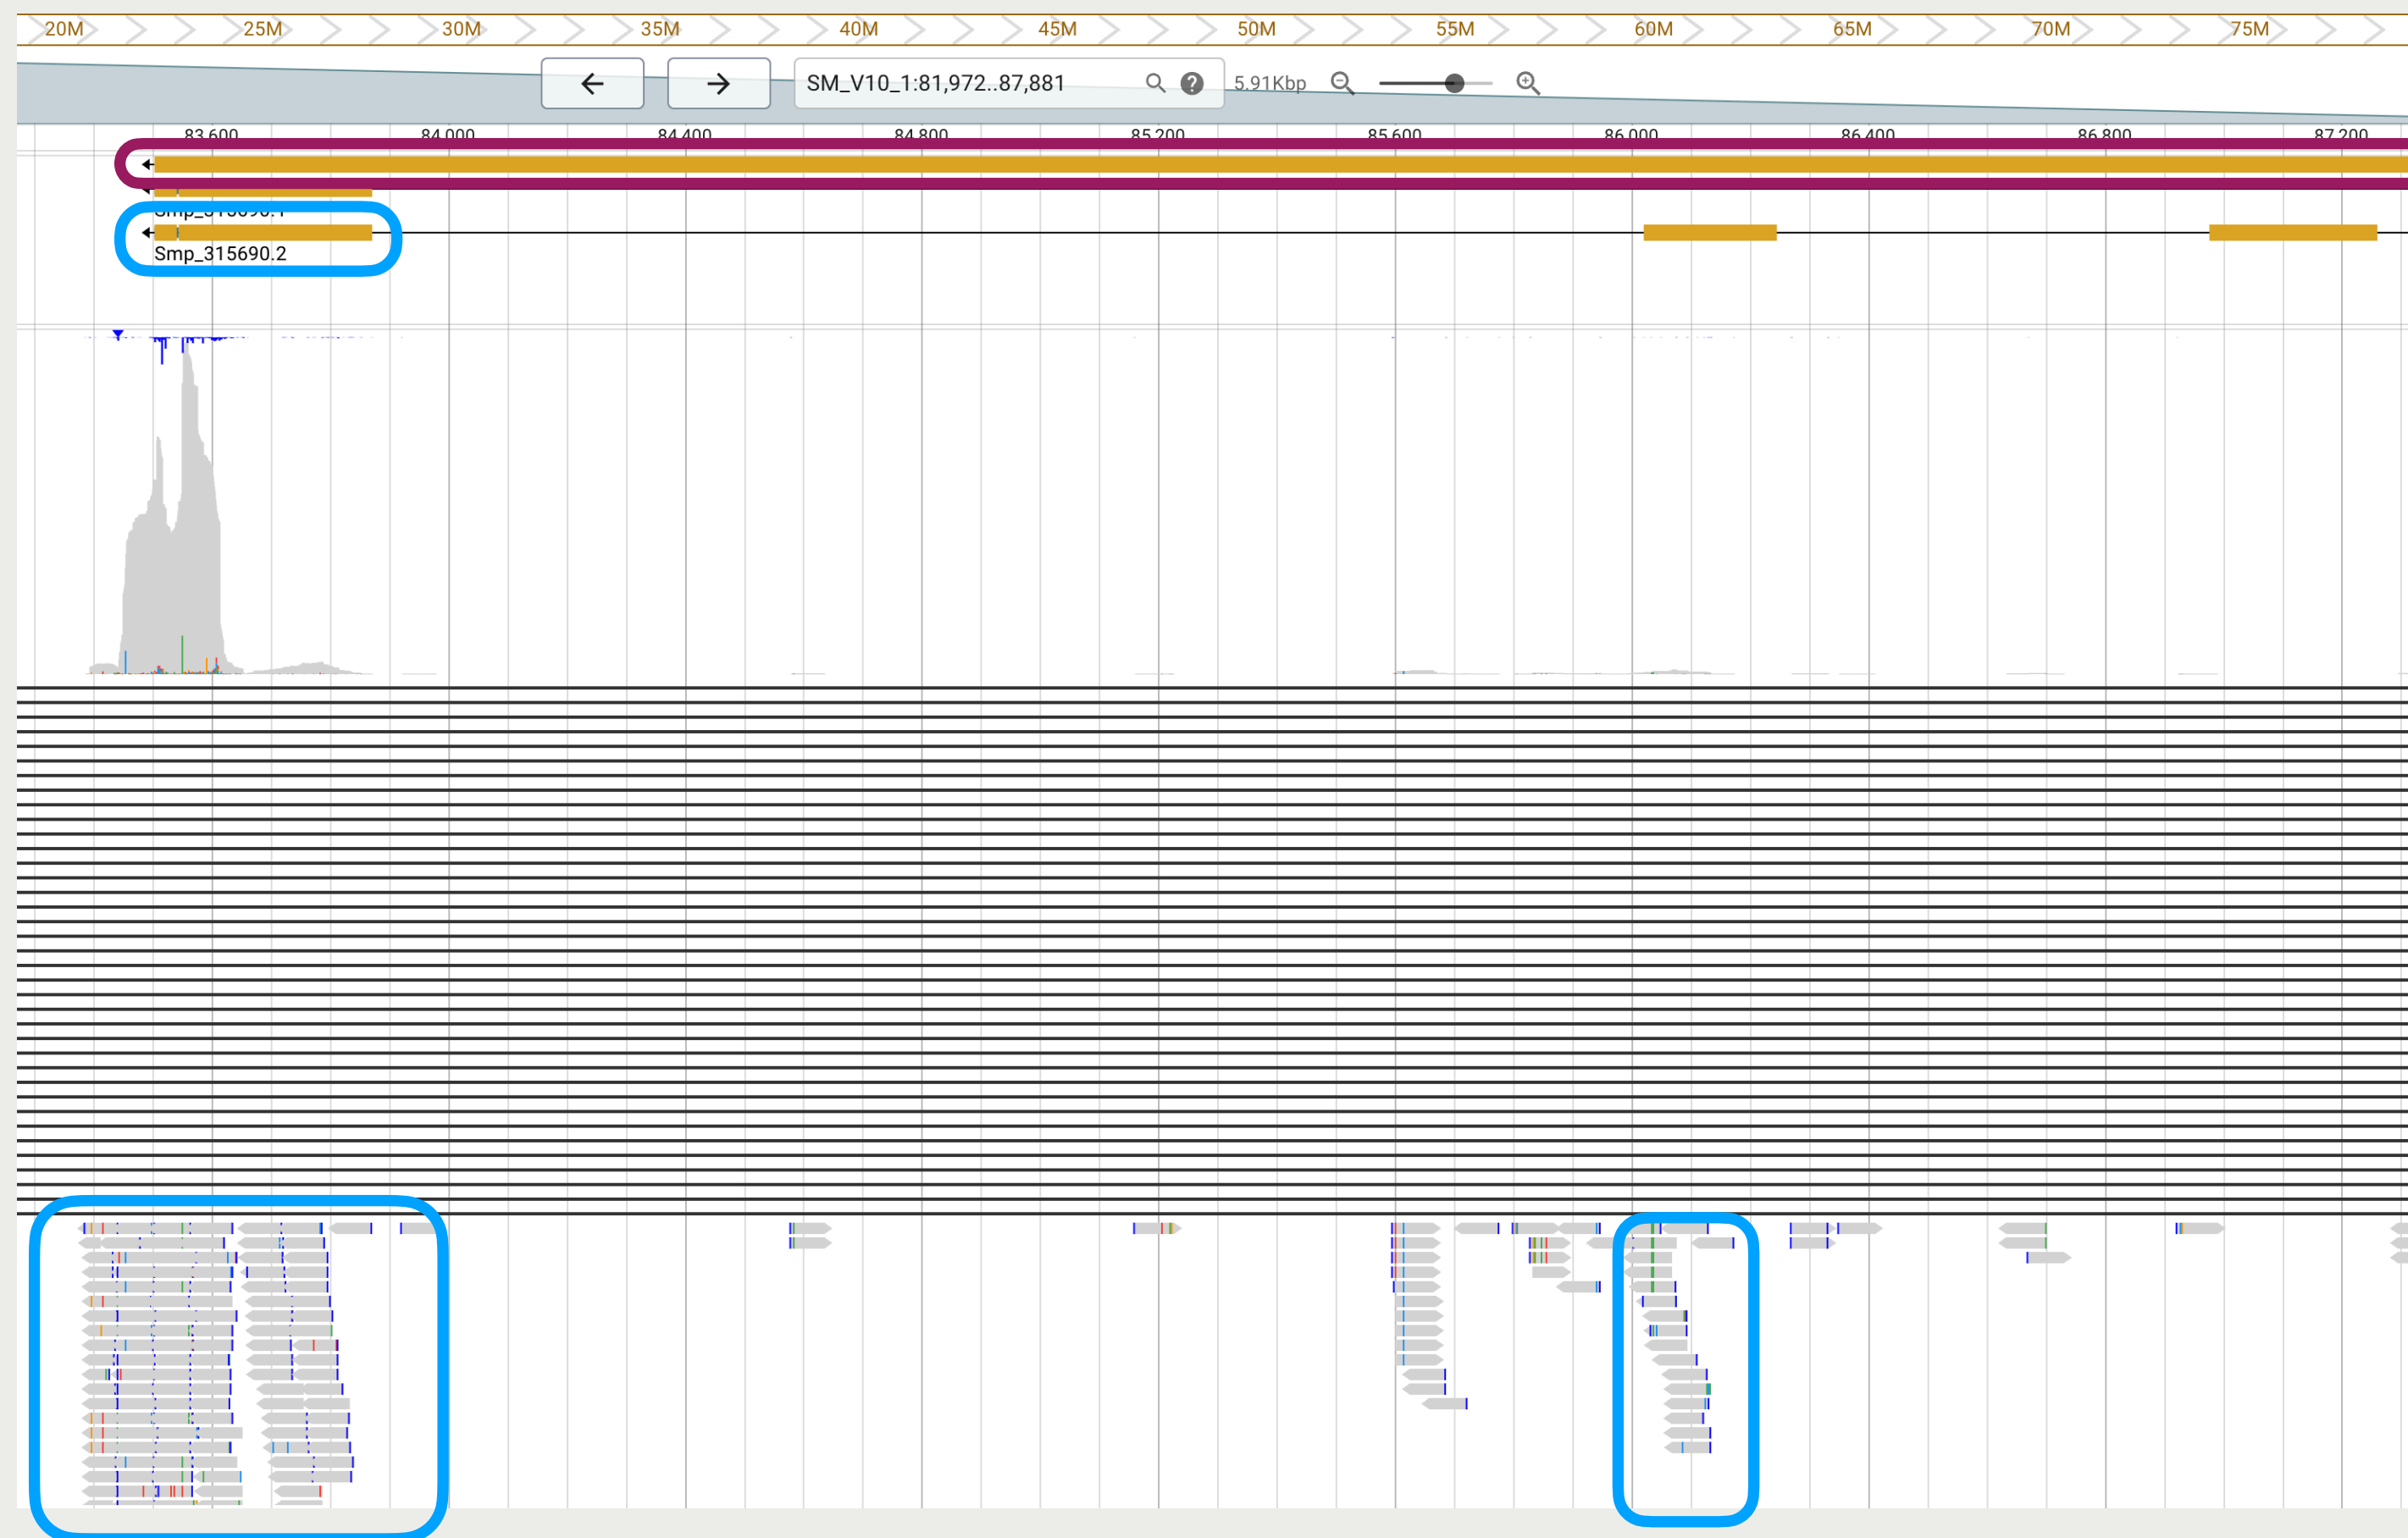
Feature = exon (by default)

Meta-feature = gene_id (by default)

- featureCounts first counts reads aligning to features (i.e., exons)
- Read is counted if it overlaps by 1 bp (by default)
- Then counts reads aligning to meta-features (i.e., genes)
 - Tallies all reads to an exon that is apart of that gene

The featureCounts approach

Features and meta-features



Hits to the feature

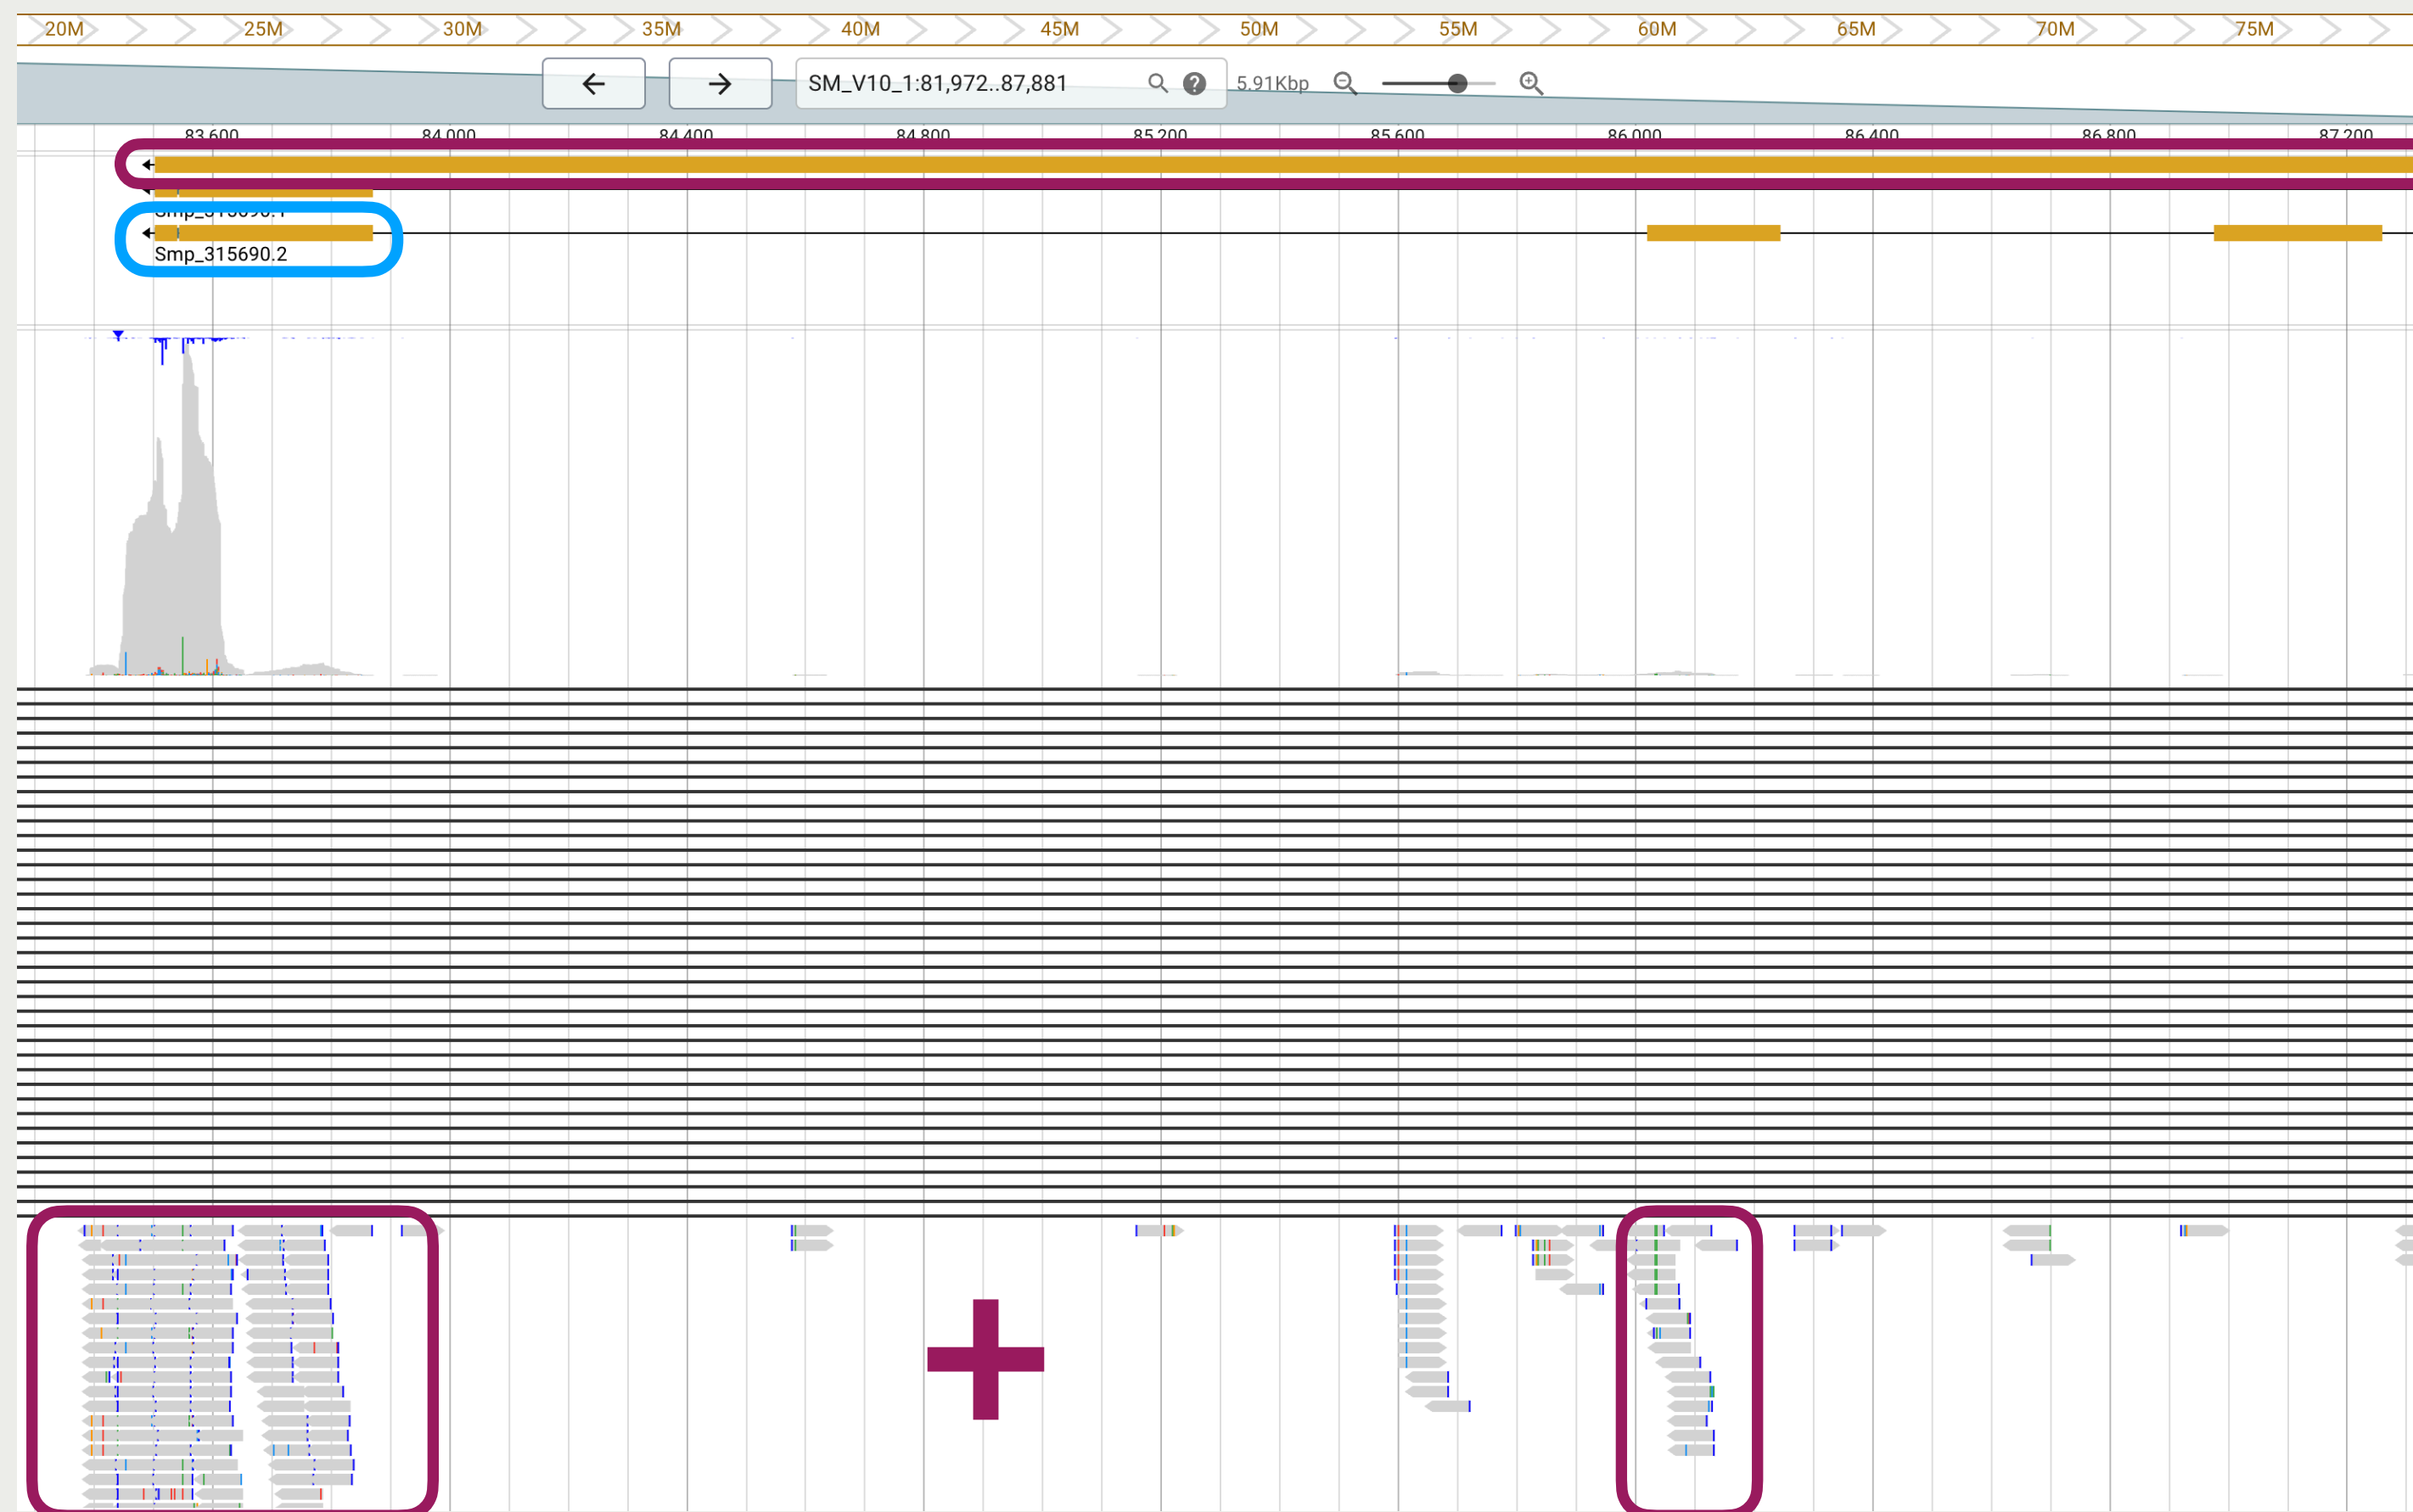
Feature = exon (by default)

Meta-feature = gene_id (by default)

- featureCounts first counts reads aligning to features (i.e., exons)
- Read is counted if it overlaps by 1 bp (by default)
- Then counts reads aligning to meta-features (i.e., genes)
 - Tallies all reads to an exon that is apart of that gene

The featureCounts approach

Features and meta-features



Feature = exon (by default)

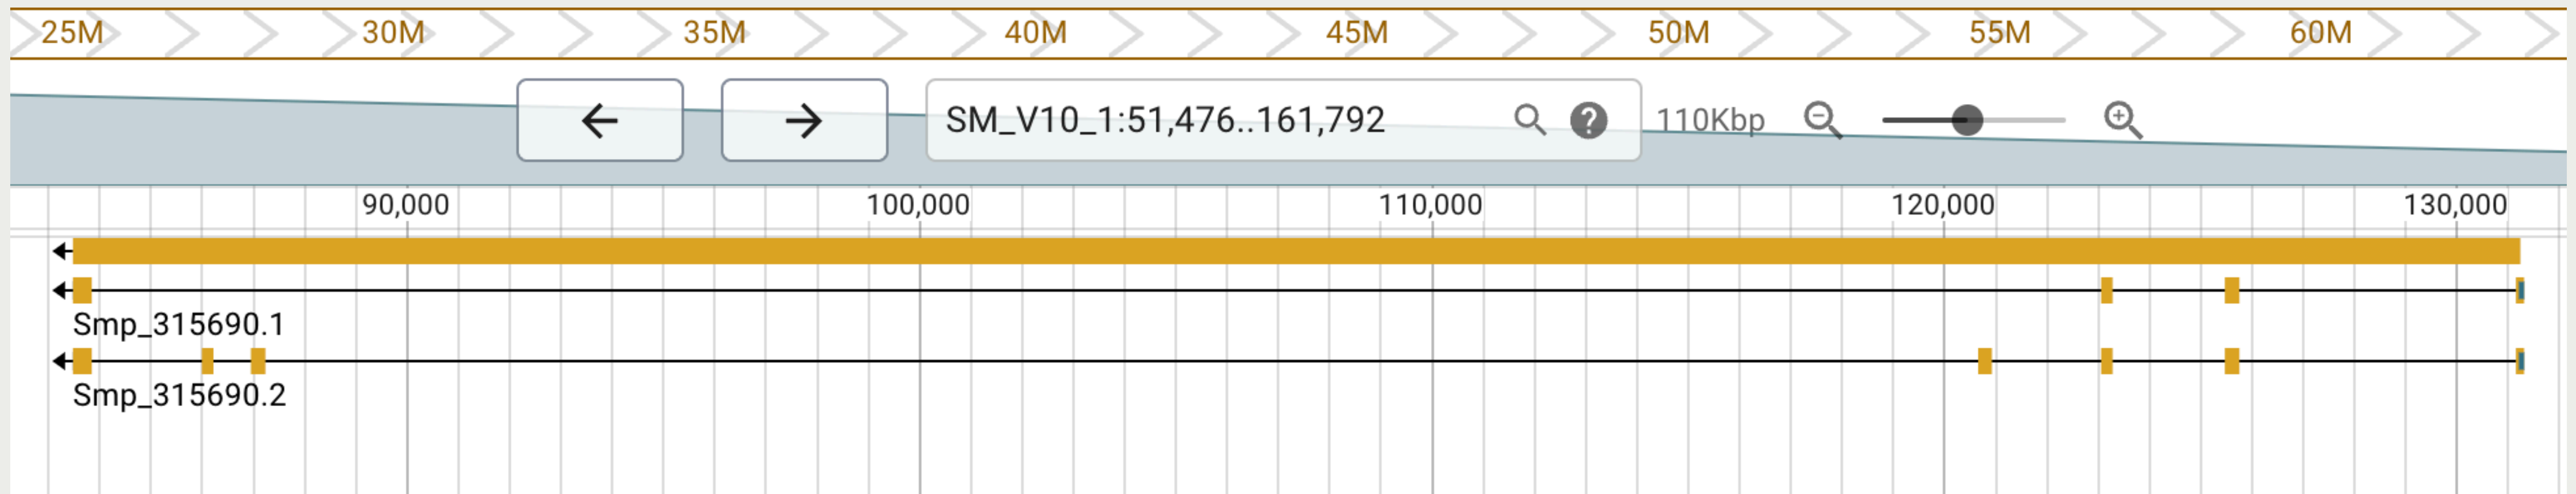
Meta-feature = gene_id (by default)

- featureCounts first counts reads aligning to features (i.e., exons)
- Read is counted if it overlaps by 1 bp (by default)
- Then counts reads aligning to meta-features (i.e., genes)
 - Tallies all reads to an exon that is apart of that gene

The featureCounts approach

Features and meta-features

- featureCounts has excellent performance (written in C)
- Works well when transcript-level differential expression is not of interest
 - What about when you want to differentiate isoforms (different transcripts from the same gene) like Smp_315690.1 and Smp_315690.2?
- Only use hits that map unambiguously...



The featureCounts approach

Features and meta-features

- Output: integer counts at meta-feature level
- Matrix of integer values - the value in the i -th row and the j -th column represents the reads assigned to gene i in sample j
- Cannot do differential expression with only these counts
 - Unnormalized
 - Longer genes will necessarily have more reads
 - Counts will vary by library size (i.e., number of reads that aligned)
 - Differential expression tools will do the normalization for us

The featureCounts approach

Features and meta-features

Geneid	Chr	Start	End	Strand	Length	ADULT_PZQ_E	ADULT_PZQ_D	ADULT_CTRL_E	ADULT_CTRL_D
Smp_329140	SM_V10_1	68427	68783	-	357	2	1	0	0
Smp_318880	SM_V10_1	261046	277059	+	16014	19	15	5	3